

Learning Precedences from Simple Symbol Features¹

Filip Bártek Martin Suda


Czech Technical University in Prague, Czech Republic

November 4, 2020

¹Supported by the ERC Consolidator grant AI4REASON no. 649043 under the EU-H2020 programme and the Czech Science Foundation project 20-06390Y.



Context



Basic concepts

Concept	Notation	Example
First-order problem	$P = (\Sigma, CI)$	
Symbols	$\Sigma = (s_1, s_2, \dots, s_n)$	
Symbol precedence	$\pi_P \sim \text{Perm}(\Sigma)$	(0, 3, 1, 2) $\begin{bmatrix} 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 \end{bmatrix}$
Order matrix	$O(\pi_P)$	
Flattened matrix	$\overrightarrow{O(\pi_P)}$	0111001000000110
Abstract solving time	$ast(P, \pi_P)$	

Main research question

Given a problem for a prover to solve,
how can we propose a good symbol precedence
using simple symbol features?

Training:

- ① For each training problem P in isolation:
 - ① Try to solve P with 100 random precedences using Vampire
 - ② Assign a cost to each of the precedences
 - ③ Assign a cost to each symbol pair by
 - ① Fitting a linear regressor of precedence cost
 - ② Extracting feature coefficients (feature \sim symbol pair)
- ② Train a symbol pair cost regressor
 - Represent each symbol by simple syntactic features

Proposing a precedence on a new problem:

- ① Predict the cost of every symbol pair
- ② Construct a precedence that minimizes the cumulative symbol pair cost

Precedence cost

For each problem $P \in \mathcal{P}_{train}$:

- ① For 100 uniformly random precedences π_P :
Run Vampire and measure $ast(P, \pi_P)$.
- ② Assign a penalty to precedences that fail to solve the problem (namely by timing out).
 - Penalty: Maximum $ast(P, \pi_P)$ over successful runs on P
- ③ Normalize the costs:
 - ① Log-scale
 - ② Standardize (ensure mean 0 and standard deviation 1)

Let $cost_{std}(\pi_P)$ be the normalized cost of π_P .

Symbol pair cost

For each problem $P \in \mathcal{P}_{train}$, train a linear model:

- Input: $O(\pi_P)$ – order matrix of a precedence

Example: $\pi_P = (0, 3, 1, 2)$, $O(\pi_P) = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 \end{bmatrix}$

- Target: $cost_{std}(\pi_P)$ – normalized cost of π_P
- Regression method: Lasso (linear regression with L_1 penalty)
- Trained model coefficients: W_P – preference matrix
 $W_{P,i,j} \sim$ cost attributed to $\llbracket s_i \text{ is ordered before } s_j \rrbracket$
- Prediction: $cost_{proxy}(\pi_P, W_P) = \overrightarrow{O(\pi_P)} \cdot \overrightarrow{W_P}$

Symbol features

6 simple features:

- Arity
- Frequency (number of occurrences)
- Number of clauses that contain the symbol
- Is it in conjecture?
- Is it in a unit clause?
- Is it introduced (Tseitin or Skolem)?

Feature vector of symbol s_i : $fv(s_i)$

Symbol pair cost model M

Generalizing across problems

- Input: feature vector $[fv(s_i), fv(s_j)]$ of a symbol pair s_i, s_j from a problem $P \in \mathcal{P}_{train}$
- Target: $W_{P,i,j}$ (cost attributed to $\llbracket s_i \text{ is ordered before } s_j \rrbracket$)
- Training sample distribution:
 - ① Sample problems from \mathcal{P}_{train} uniformly.
 - ② Given a problem, weight symbol pairs by absolute target value.
- Learning method:
 - Elastic-Net (linear regression with L_1 and L_2 penalties)
 - Gradient Boosting Machine
- Prediction: $M([fv(s_i), fv(s_j)])$

Proposing a good precedence

Given a previously unseen problem P :

- ① Estimate a preference matrix \widehat{W}_P .
$$\widehat{W}_{P,i,j} = M([fv(s_i), fv(s_j)]).$$
- ② Construct a precedence $\widehat{\pi}_P$ that approximately minimizes
$$cost_{proxy}(\widehat{\pi}_P, \widehat{W}_P) = \overrightarrow{O(\widehat{\pi}_P)} \cdot \overrightarrow{\widehat{W}_P}.$$
 - NP-hard in general
 - 2-approximation greedy algorithm by Cohen et al. [1]:
Repeatedly append symbol with the smallest potential:
$$c(s_i) = \sum_{s_j \in \Sigma_{\text{avail}}} \widehat{W}_{P,i,j} - \sum_{s_j \in \Sigma_{\text{avail}}} \widehat{W}_{P,j,i}$$

Experimental evaluation

- Learning only predicate precedences
- 5 evaluation iterations, each using:
 - 1000 training problems
 - 1000 test problems

Case	Successes out of 1000	
	Mean	Std
Best of 10 random	501.8	13.85
invfreq	475.8	13.83
Elastic-Net	471.6	10.21
Gradient Boosting	464.8	13.78
Elastic-Net without weighting	461.2	8.11
Random	450.4	10.25

Elastic-Net feature coefficients of symbol pairs

Split	Left symbol			Right symbol		
	Arity	Freq.	Unit freq.	Arity	Freq.	Unit freq.
0		-0.01		-0.98	0.01	
1		-0.48			0.08	0.44
2			-0.64		0.36	
3	0.88	-0.03	0.01		-0.03	0.05
4		-0.62			0.30	0.07
\mathcal{P}_{train}			-0.57		0.43	

The other 3 features have all zero coefficients.

- Training sets 1, 2, 4, \mathcal{P}_{train} :
Similar to `invfreq` heuristic (low frequency \sim early inference)
- Training sets 0, 3:
Similar to arity heuristic (high arity \sim early inference)

Summary

We designed a system that:

- Trains on random precedences using preference matrices
- Proposes a good precedence using a greedy algorithm

We evaluated the system on TPTP for predicate precedences.

Future:

- Stabilize training
- Evaluate also for function symbol precedences
- Richer symbol representation using graph neural networks

Section 5

Bonus slides

References

- [1] William W. Cohen, Robert E. Schapire, and Yoram Singer. Learning to order things. *CoRR*, abs/1105.5464, 2011. URL <http://arxiv.org/abs/1105.5464>.
- [2] G. Sutcliffe. The TPTP Problem Library and Associated Infrastructure. From CNF to TH0, TPTP v6.4.0. *Journal of Automated Reasoning*, 59(4):483–502, 2017.

Experimental setup

- Only predicate precedences are learned.
Function symbols are ordered by `invfreq`.
- Problems from TPTP [2] – CNF and FOF (clausified with Vampire)
 - \mathcal{P}_{train} (8217 problems): at most 200 predicate symbols, at least 1 out of 24 random predicate precedences yield success
 - \mathcal{P}_{test} (15751 problems): at most 1024 predicate symbols
- 5 evaluation iterations (splits): 1000 training problems and 1000 test problems
- 100 precedences per training problem
- Vampire configuration: time limit: 10 seconds, memory limit: 8192 MB, literal comparison mode: predicate, function symbol precedence: `invfreq`, saturation algorithm: discount, age-weight ratio: 1:10, AVATAR: disabled
- 10^6 symbol pair samples to train M

Elastic-Net feature coefficients of individual symbols

Training set	Arity	Frequency	Unit frequency
0	-.98	.01	-.01
1		.56	.44
2		.36	.64
3	-.88		.04
4		.93	.07
\mathcal{P}_{train}		.43	.57

Symbol order: descending by predicted value

- Sets 1, 2, 4, \mathcal{P}_{train} :
 - Descending by frequency: low frequency \sim early inference
 - Similar to `invfreq` and `vampire --sp frequency`
- Sets 0, 3:
 - Ascending by arity: high arity \sim early inference
 - Similar to `vampire --sp arity`